

Data Preparation

Paul Rodriguez



On the Importance of Data Prep

- **“Garbage in, garbage out”**
- **Sometimes takes 60-80% of the whole data mining effort**


Working definition

- **Data Preparation:**
 - Cleaning
 - Filtering
 - Transforming
 - Organizing the data matrix (aka 'data wrangling' or 'data munging')
 - Variable Selection/Dimension Reduction

In a nutshell, know your variables and prepare data for modeling

Statistical Noise:

- Outliers
e.g. remove them,




mean + 3*std-dev

Statistical Noise:

- Outliers

e.g. remove them, but cutoff sometimes seems arbitrary



mean + 3*std-dev

Cleaning Noise

- Entity Resolution and Record Linkage

e.g. Are these equal?

West Main Street

W Main St

Strategy:

use dictionaries and search possible matches

Missing Data (NA vs NULL)

- **Not applicable - NA**

e.g. spouse name depends on marital status

- **Not available - NULL**

unknown

not entered

- **In R:**

NA is logically missing (ie the value is not available – null)

NULL is object not yet defined (ie the object is not available)

Missing Data

- **What are frequency counts of missing variables?**

Are entries missing completely at random or contingent on some other variable?

Quick Approaches

- Delete instances
and/or
- Delete attributes with high missing-ness

Quick Approaches

- Leave as missing
 - Some algorithms implementation handle missing values (ie Decision Trees)

Simple Imputation

- Use the attribute mean
- Use the attribute mean for each class label

Complicated Imputation

- Use a model (based on other attributes) to infer missing value

Complicated Imputation

- Use a model (based on other attributes) to infer missing value

*Best strategy depends on
time vs accuracy tradeoffs*

R and missing data

- Several packages, such as 'mice', 'amelia'
- Produces multiple data sets
- Iterates over missing data estimates and linear model estimates

Mice uses Gibbs sampling (slower)

Amelia uses Expectation Maximization (faster)

- **Beware of correlation in variables**

Matrices not invertible

R and missing data

- ‘Amelia’ package example
 - 50 attributes from UN voting data
 - 1K-100K entries missing per col for about 20 cols
 - 300K rows ~ 1 hour on Gordon compute node (not run on the user’s PC)

```
# run the imputation
library('amelia')
a.out <- amelia(data, ts = "year", cs = "dyadid",
               idvars = c("dyadidyr", "cntryera", "statea", "stateb"),
               intercs=FALSE, p2s = 2, m=10, parallel = "multicore")
```

*time variable for time
series models*

*cross section to select
temporal periods*

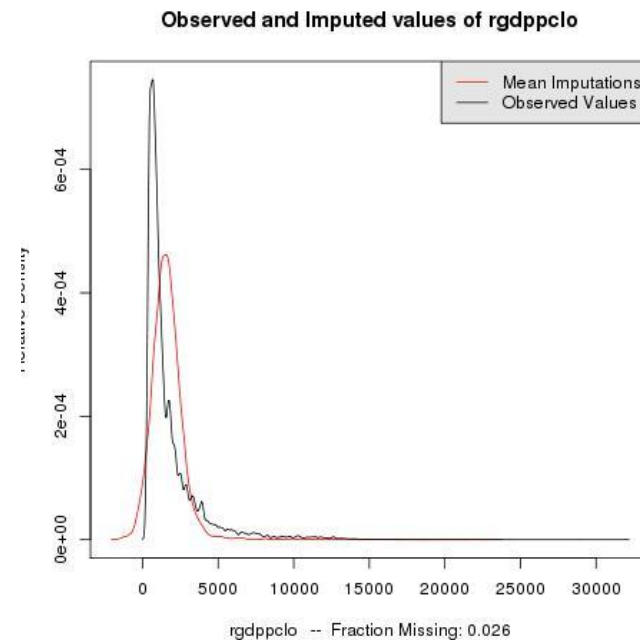
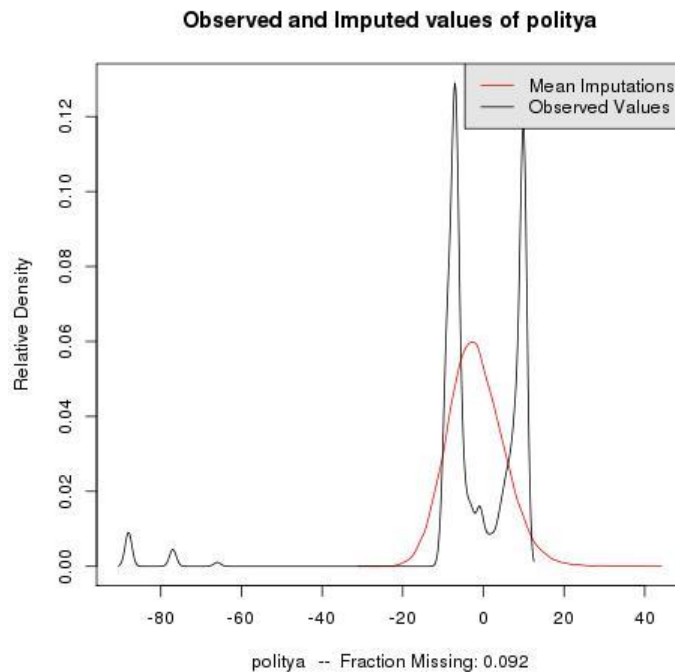
ignore 'id' variables

interactions

parallel options

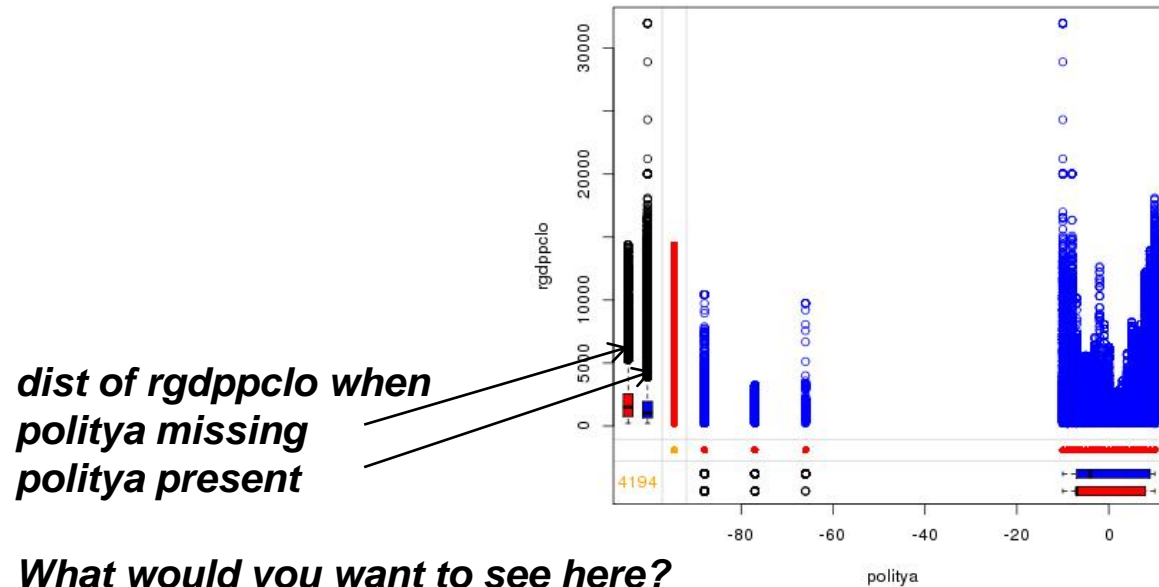
R and missing data

#QA on missing data by comparing density of imputed & original data
`compare.density(a.out, var="politya")`
`compare.density(a.out, var='rgdppcontg')`



R and missing data

```
# Useful library for printing margin plots, to compare histograms  
# conditional on missing/non-missing data  
library('VIM')  
marginplot(gart2use[,c('politya', 'rgdppclo')],  
           col=c('blue','red','orange'))
```



Variable Transformations

- **Engineer new features**
- **Combine attributes**
e.g. rates and ratios
- **Normalize or Scale data**
- **Discretize data**
(perhaps more intuitive to deal with binned data)

Feature Engineering is Variable Enhancement

- Use Domain and world knowledge
- **Example: given date and location of doctor visits**
 - a new variable for Number-of-1st-time-visits
 - deduce a new variable for Number-of-visits-over-25-miles
 - deduce a new variable for Amount-of-time-between-visits

Re-scaling

- **Mean center** $x_{new} = x - \text{mean}(x)$
- **z-score** $score = \frac{x - \text{mean}(x)}{\text{std}(x)}$
- **Scale to [0...1]** $x_{new} = \frac{x - \min(x)}{\max(x) - \min(x)}$
- **log scaling** $x_{new} = \log(x)$

Generally

- **Preparing data is based on statistical principles,**
- **But also heuristics**

Data Wrangling Exercise: Weather Data

weather - Excel

FILE HOME INSERT PAGE LAYOUT FORMULAS DATA REVIEW VIEW ACROBAT

Clipboard Font Alignment Number Styles Cells Editing

Calibri 18 A A Wrap Text Merge & Center General \$ % .00 0.00 Conditional Formatting Format as Table Cell Styles Insert Delete Format AutoSum Fill Clear Sort & Find & Filter Select

A1 X ✓ fx Date

	A	B	C	D	E	F	G	H	I	J
1	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDi	WindGustSp	WindDir9am
2	11/1/2007	Canberra	8	24.3	0	3.4	6.3	NW	30	SW
3	11/2/2007	Canberra	14	26.9	3.6	4.4	9.7	ENE	39	E
4	11/3/2007	Canberra	13.7	23.4	3.6	5.8	3.3	NW	85	N
5	11/4/2007	Canberra	13.3	15.5	39.8	7.2	9.1	NW	54	WNW
6	11/5/2007	Canberra	7.6	16.1	2.8	5.6	10.6	SSE	50	SSE
7	11/6/2007	Canberra	6.2	16.9	0	5.8	8.2	SE	44	SE
8	11/7/2007	Canberra	6.1	18.2	0.2	4.2	8.4	SE	43	SE
9	11/8/2007	Canberra	8.3	17	0	5.6	4.6	E	41	SE
10	11/9/2007	Canberra	8.8	19.5	0	4	4.1	S	48	E
11	11/10/2007	Canberra	8.4	22.8	16.2	5.4	7.7	E	31	S
12	11/11/2007	Canberra	9.1	25.2	0	4.2	11.9	N	30	SE
13	11/12/2007	Canberra	8.5	27.3	0.2	7.2	12.5	E	41	E
14	11/13/2007	Canberra	10.1	27.9	0	7.2	13	WNW	30	S

weather

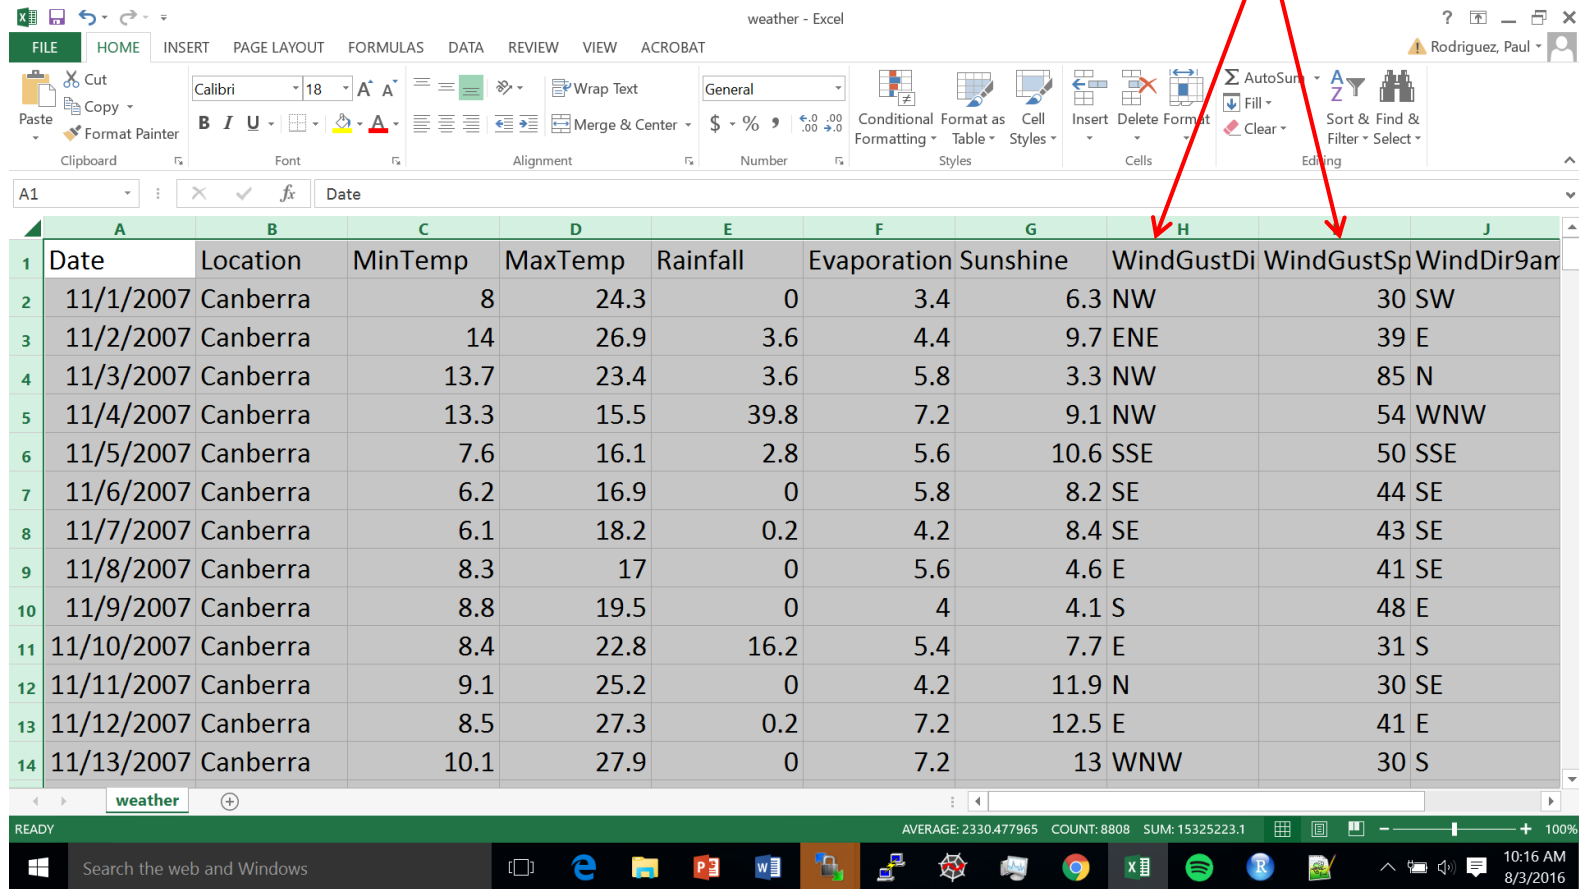
READY AVERAGE: 2330.477965 COUNT: 8808 SUM: 15325223.1 100%

Search the web and Windows

10:16 AM 8/3/2016

Transforming Weather Data Matrix

*Let's consider WindGustDir as a repeated measurement
Do we want that all in one row? Or in their own row?
- depends on the analysis*



	A	B	C	D	E	F	G	H	I	J
1	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDi	WindGustSp	WindDir9arr
2	11/1/2007	Canberra	8	24.3	0	3.4	6.3	NW	30	SW
3	11/2/2007	Canberra	14	26.9	3.6	4.4	9.7	ENE	39	E
4	11/3/2007	Canberra	13.7	23.4	3.6	5.8	3.3	NW	85	N
5	11/4/2007	Canberra	13.3	15.5	39.8	7.2	9.1	NW	54	WNW
6	11/5/2007	Canberra	7.6	16.1	2.8	5.6	10.6	SSE	50	SSE
7	11/6/2007	Canberra	6.2	16.9	0	5.8	8.2	SE	44	SE
8	11/7/2007	Canberra	6.1	18.2	0.2	4.2	8.4	SE	43	SE
9	11/8/2007	Canberra	8.3	17	0	5.6	4.6	E	41	SE
10	11/9/2007	Canberra	8.8	19.5	0	4	4.1	S	48	E
11	11/10/2007	Canberra	8.4	22.8	16.2	5.4	7.7	E	31	S
12	11/11/2007	Canberra	9.1	25.2	0	4.2	11.9	N	30	SE
13	11/12/2007	Canberra	8.5	27.3	0.2	7.2	12.5	E	41	E
14	11/13/2007	Canberra	10.1	27.9	0	7.2	13	WNW	30	S

Data Wrangling exercise

#Read in data:

```
W_df = read.table('weather.csv',  
                  header=TRUE,  
                  sep=";",  
                  stringsAsFactors = TRUE)
```


Long to Wide transform

	A	B	C	D	E	F	G	H	I	J
1	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDi	WindGustSp	WindDir9am
2	11/1/2007	Canberra	8	24.3	0	3.4	6.3	NW	30	SW
3	11/2/2007	Canberra	14	26.9	3.6	4.4	9.7	ENE	39	E
4	11/3/2007	Canberra	13.7	23.4	3.6	5.8	3.3	NW	85	N

date, location and the rest identify the row

WindGustDir entries are labels for the repeated measures

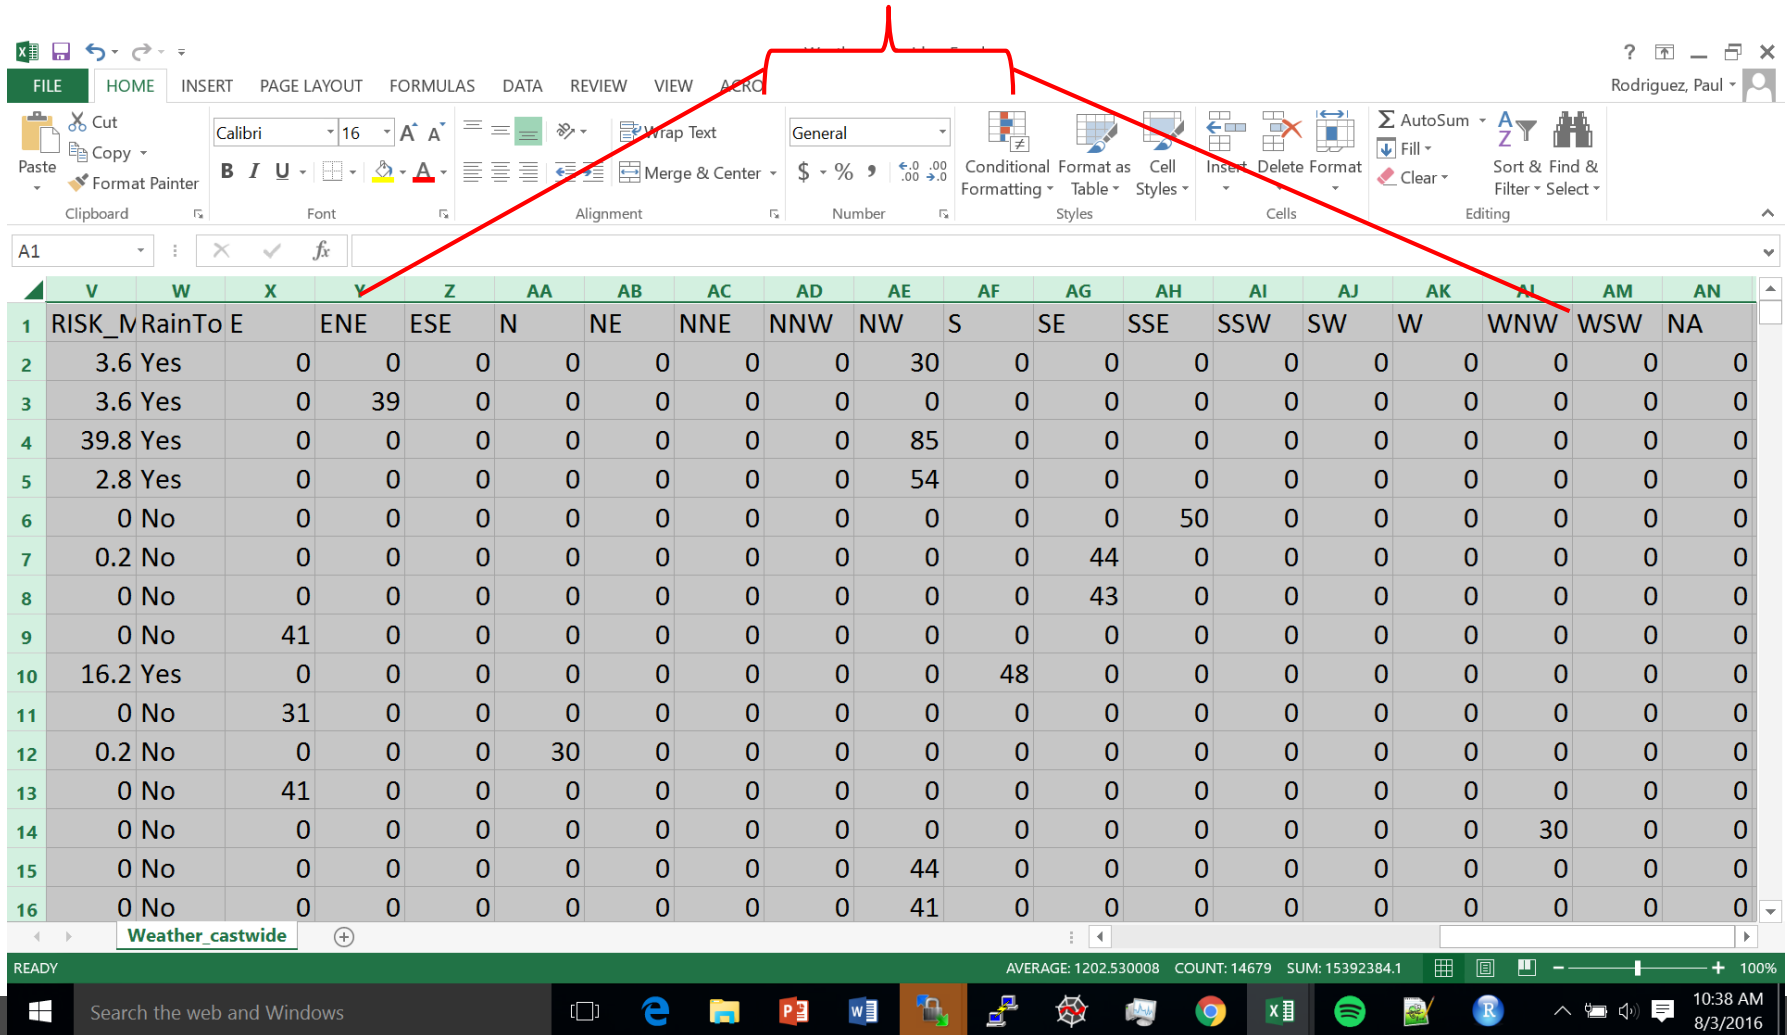
```
install.packages("reshape2")
library(reshape2)
W_wide = dcast(W_df,
               formula = Date + Location + ... ~ WindGustDir,
               fill = 0,
               value.var = "WindGustSpeed")
```

this could be 0 or NA

indicate variable that has the repeated measurement values

Transformed Data Matrix

Now: WindGustDir values each have their own column



	V	W	X	Y	Z	AA	AB	AC	AD	AE	AF	AG	AH	AI	AJ	AK	AL	AM	AN
1	RISK_N	RainTo	E	ENE	ESE	N	NE	NNE	NNW	NW	S	SE	SSE	SSW	SW	W	WNW	WSW	NA
2	3.6	Yes		0	0	0	0	0	0	0	30	0	0	0	0	0	0	0	0
3	3.6	Yes		0	39	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	39.8	Yes		0	0	0	0	0	0	85	0	0	0	0	0	0	0	0	0
5	2.8	Yes		0	0	0	0	0	0	54	0	0	0	0	0	0	0	0	0
6	0	No		0	0	0	0	0	0	0	0	0	50	0	0	0	0	0	0
7	0.2	No		0	0	0	0	0	0	0	0	44	0	0	0	0	0	0	0
8	0	No		0	0	0	0	0	0	0	0	43	0	0	0	0	0	0	0
9	0	No	41	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
10	16.2	Yes		0	0	0	0	0	0	0	48	0	0	0	0	0	0	0	0
11	0	No	31	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
12	0.2	No		0	0	0	30	0	0	0	0	0	0	0	0	0	0	0	0
13	0	No	41	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
14	0	No		0	0	0	0	0	0	0	0	0	0	0	0	0	30	0	0
15	0	No		0	0	0	0	0	0	44	0	0	0	0	0	0	0	0	0
16	0	No		0	0	0	0	0	0	41	0	0	0	0	0	0	0	0	0

Extra to try:

```
# wide to long: ie 'melt' repeated measure into long  
# table
```

```
W_melt =melt(W_cast,  
             na.rm=TRUE,  
             measure.vars=c(23:39),  
             variable.name="WindGustDir_cast")
```

*The repeated measures are
now in columns 23 to 30*

pause

Reading Material

- **Data Preparation for Data Mining by Dorian Pyle**
 - http://www.ebook3000.com/Data-Preparation-for-Data-Mining_88909.html
- **Data mining – Practical Machine learning tools and techniques by Witten & Frank**
 - <http://books.google.com>

Many Variables

- **More variables \Rightarrow more information, but also more noise and more ways of interactions**
- **2 ways to handle many variables**
 - Variable Selection
 - Dimension reduction methods

Variable selection vs Dimensionality Reduction

- **Prior to algorithm, depends on data**
 - For large P , with noise particular to variables, try variable selection
 - For large P , diffuse noise, try dimension reduction by Matrix Factorization

Variable selection

- **Heuristic methods:**
 - remove variables with low correlations to outcome
 - (other criteria: information gain, sensitivity, etc...)
- **Step wise: add 1 variable at a time and test algorithm on samples**

Variable selection

- **Some algorithms are robust to extra noise variables**
- **E.g. Least Angle Regression (L_1 penalty),**
penalize small effect sizes (zero them out)
- **E.g. Random Forest outputs ‘importance’**
low importance implies non-influence in the model
(other criteria: information gain, sensitivity, etc...)

Matrix Factorization:

Given a numeric matrix, can we reduce the number of columns?

conversely

Can we find interesting subspaces?

Matrix Factorization:

Given a numeric matrix, can we reduce the number of columns?

- Yes, if features are constant or redundant

Matrix Factorization:


Given a numeric matrix, can we reduce the number of columns?

- Yes, if features are constant or redundant
- Yes, if features only contribute noise
(conversely, want features that contribute to variations of the data)

Example: 2D data


Weight Height

S1	50	179
S2	66	175
S3	74	180
S4	94	192




***this is
the
input
space***

Example: Athletes' Height by Weight




Find a line that aligns with the data.

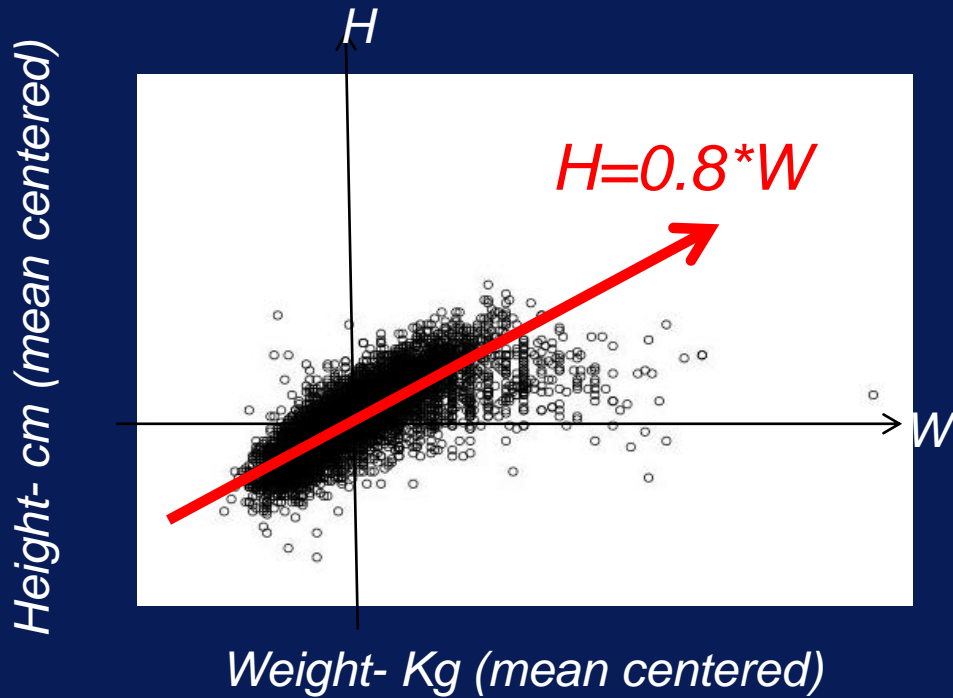
Example: Athletes' Height by Weight




Find a line that aligns with the data.



Find a line that aligns with the data.




Find a line that aligns with the data.



Note that $W=1, H=0.8$ is a point on the line, for example.


Find a line that aligns with the data.




Note that $W=1, H=0.8$ is a point on the line, for example.

Let $[1 \ 0.8]$ represent the line, as a combination of W & H .

Find a line that aligns with the data.




The next direction of most variance.



Now Rotate Axis

*New axis (AKA features)
defined as combinations of
old features*

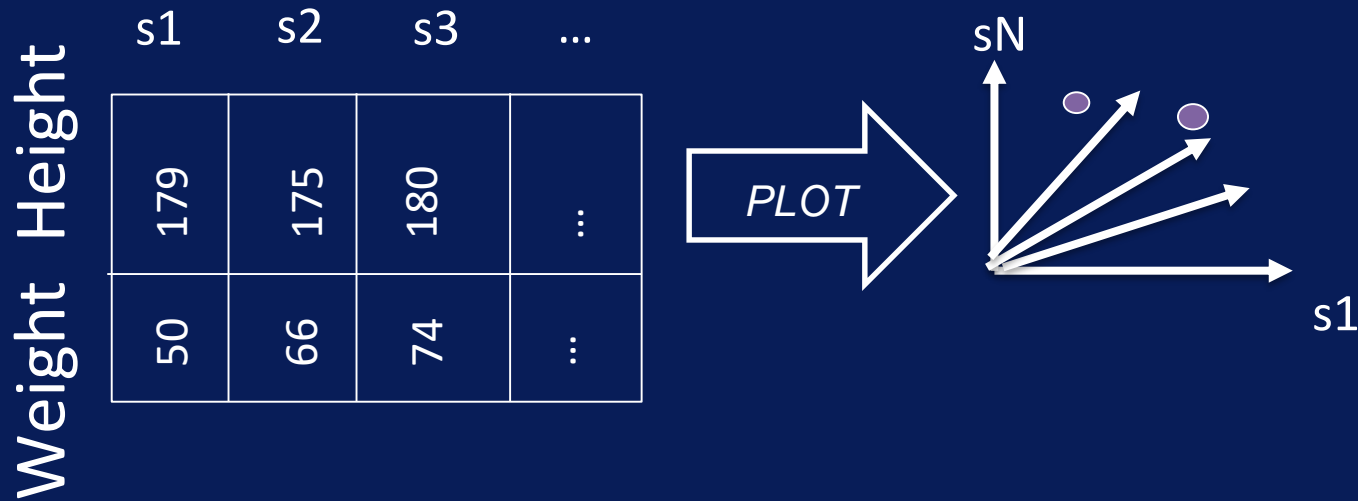
H *Possibly reduce dimensions*



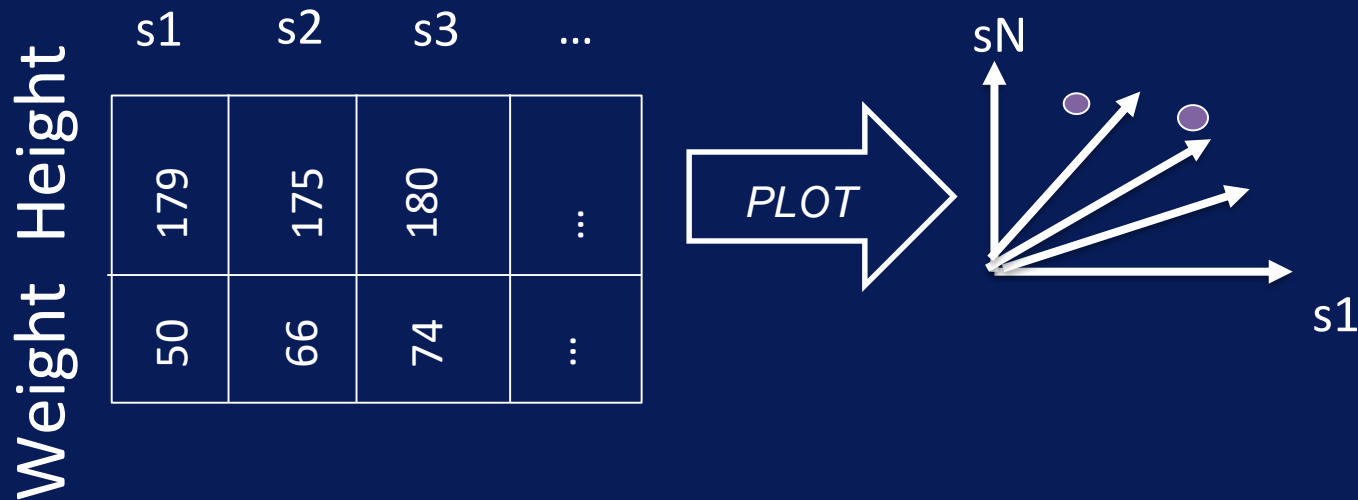
Project all points to one axis

(defined by the $[1 \ 0.8]$ 2D vector)

2D data transposed to 2 points in high dimensional space




2D data transposed to 2 points in high dimensional space



Same process as before, but now factors are N-dimensional vectors

Any Matrix can be approximated by outer product of factors



Original matrix




Factor A




Factor B


Any Matrix can be approximated by outer product of factors



Any Matrix can be approximated by outer product of factors



More factors gives better approximation




- Best Known Factorization Algorithms:
SVD (singular value decomposition)
PCA (principle component analysis)

- Best Known Factorization Algorithms:
SVD (singular value decomposition)
PCA (principle component analysis)

*Find orthogonal factors and
scale them down (i.e. normalize)*

SVD: factors and 'singular' scale values



- More generally:

Factorization Algorithms may vary depending on criterion for how factors 'align' with data.


- More generally:

Factorization Algorithms may vary depending on criterion for how factors 'align' with data.

- Number of factors to use depends on tradeoff of good approximation vs good dimensional reduction

Can use cross validation or heuristics to choose.

Note: PCA conserves and reorders variance



Total Variance
Conserved:
 Var in Weight
+
 Var in Height
=
 $\text{Var in PC1} +$
 Var in PC2

In general:
 $\text{Var in PC1} >$
 $\text{Var in PC2} >$
 $\text{Var in PC3} \dots$

Summary: Principle Components

- Can choose k heuristically as approximation improves, or choose k so that high percent (ie 80-95%) of data variance accounted for
- aka Singular Value Decomposition
 - PCA on square matrices only
 - SVD gives same vectors on square matrices
- Works for numeric data only
- For higher dimensional data, use factors to visualize 2 factors at a time

SVD Exercise

- Overview

Run on numeric fields of weather data

Run SVD and select smaller number of dimensions

Run linear model with original and reduced data

Later, we'll compare SVD components with Clustering

#W_num is only numeric or integer fields of Weather data

```
> Wsvd=svd(W_num)
```

```
> str(Wsvd)
```


List of 3

```
$ d: num [1:9] 27442.7 231.2 96.4 68.2 44.5 ...
```

```
$ u: num [1:363, 1:9] -0.0524 -0.0521 -0.052 -0.0519 -0.0525 ...
```

```
$ v: num [1:9, 1:9] -0.005042 -0.014276 -0.000969 -0.00314 -0.005491 ...
```

Exercise highlights



Compare Linear Model results, using $Y = \text{raintomorrow}$:

**look for residual standard error values and degree of freedom,
look at coefficient estimates**

Call:

lm(formula = Ymc ~ W_mncntr)

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.124e-15	1.641e-02	0.000	1.000000
W_mncntrMinTemp	-1.368e-02	1.013e-02	-1.350	0.177844
W_mncntrMaxTemp	1.035e-02	2.010e-02	0.515	0.607120
W_mncntrRainfall	4.269e-03	4.471e-03	0.955	0.340442
W_mncntrEvaporation	2.690e-02	1.010e-02	2.663	0.008137 **
W_mncntrSunshine	-3.446e-02	9.898e-03	-3.482	0.000570 ***
W_mncntrPressure9am	6.569e-02	1.325e-02	4.960	1.16e-06 ***
W_mncntrPressure3pm	-8.047e-02	1.337e-02	-6.021	4.89e-09 ***

Residual standard error: 0.2971 on 311 degrees of freedom

Call:

lm(formula = Ymc ~ W_dfred)

Coefficients: (13 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-4.874e-16	1.808e-02	0.000	1.000000
W_dfred1	4.519e+00	1.242e+00	3.638	0.000320 ***
W_dfred2	4.650e+00	1.307e+00	3.559	0.000429 ***
W_dfred3	1.580e+00	4.357e-01	3.627	0.000333 ***
W_dfred4	NA	NA	NA	NA
W_dfred5	NA	NA	NA	NA

Residual standard error: 0.3274 on 324 degrees of freedom

- end