

# 2017 SDSC Summer Institute Scalable Machine Learning



# Scalable Machine Learning Agenda

**1:30 - 2:15 – R in HPC**

**2:15 - 3:15 – Machine Learning with Spark**

**3:15 - 3:30 – Break**

**3:30 - 4:15 – Deep Learning Overview**

**4:15 - 4:45 – CNN Transfer Learning with Keras**

**4:45 - 5:00 – Wrap-up**

# Machine Learning with Spark

Mai H. Nguyen

# Spark Topics

- **Spark Overview**
- **Programming in Spark**
- **MLlib**

# Spark Overview

# What is Spark?



- **General framework for distributed computing**
- **Provides built-in data parallelism and fault-tolerance for big data processing on a cluster**
- **Goals: speed, ease of use, generality**
  - Multiple analytics applications, data sources, platforms
- **Open-source**

# Basics of Distributed Processing with Spark

Expressive programming environment

In-memory processing

Support for diverse workloads

Interactive shell

# The Spark Stack



**SparkSQL**

**Spark  
Streaming**

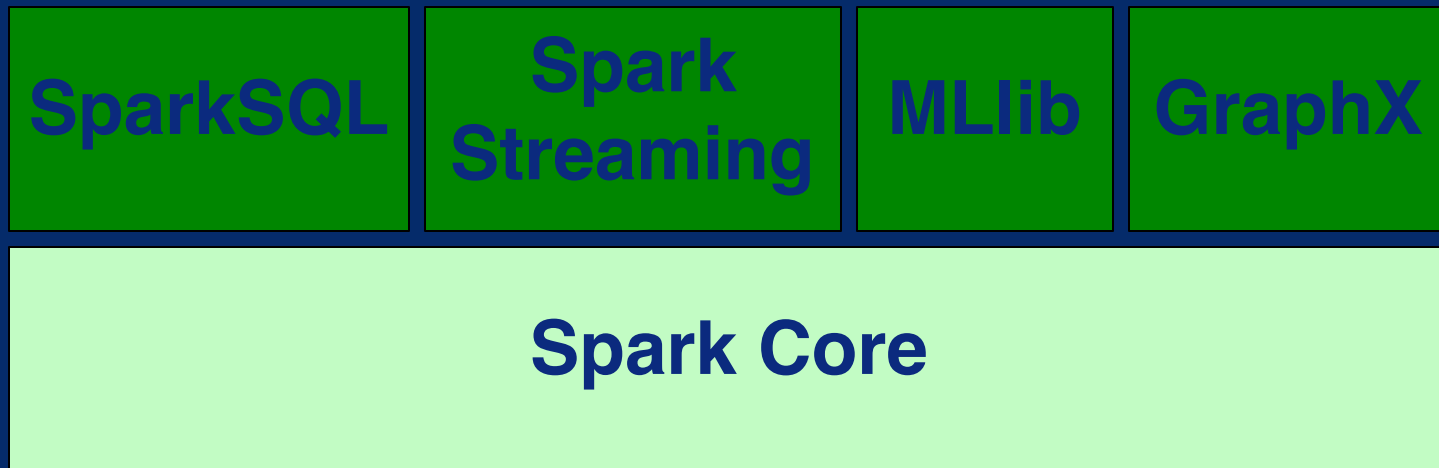
**MLlib**

**GraphX**

**Spark Core**

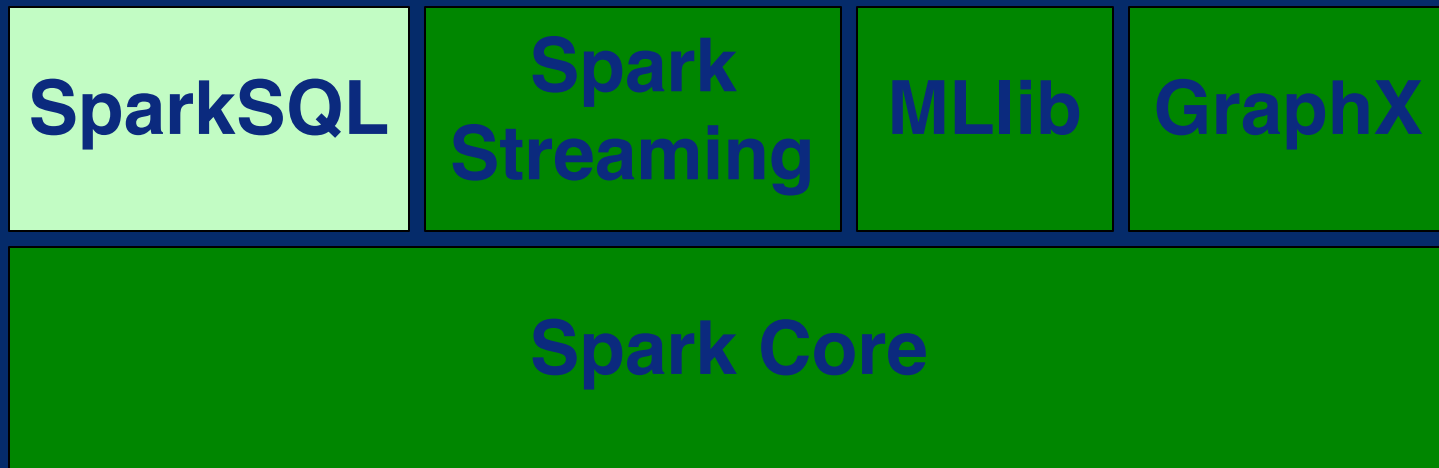


# The Spark Stack



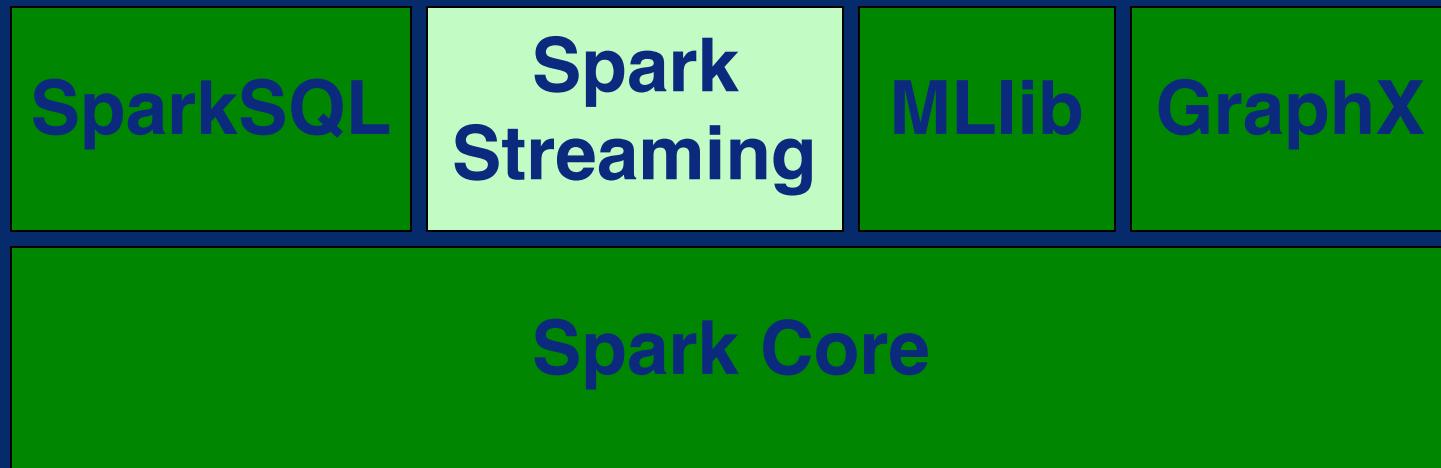
**Distributed computing**

# The Spark Stack



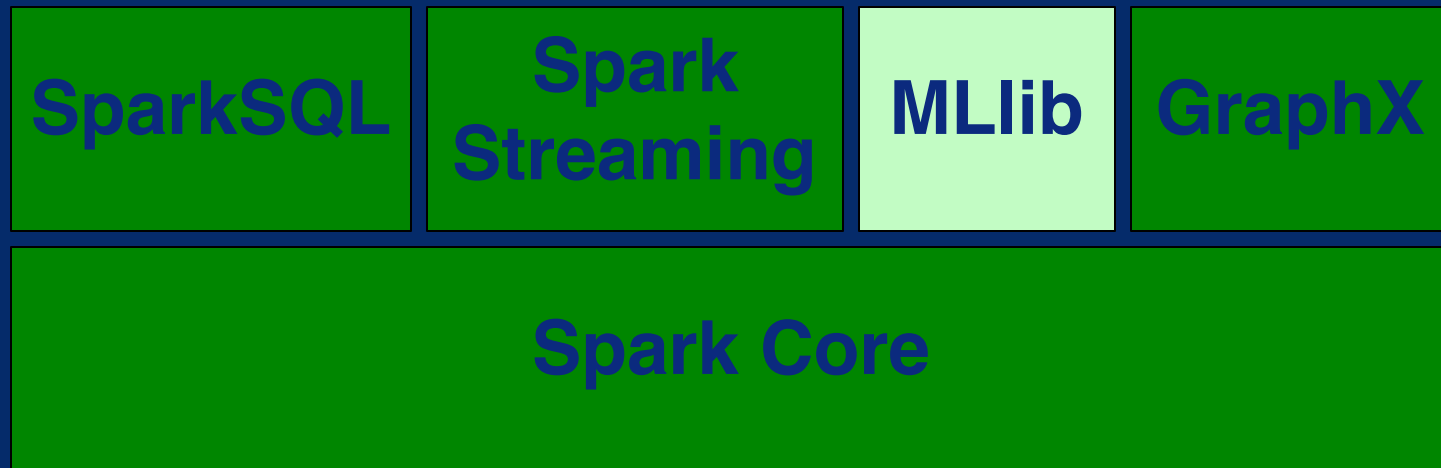
**SQL-like querying**

# The Spark Stack



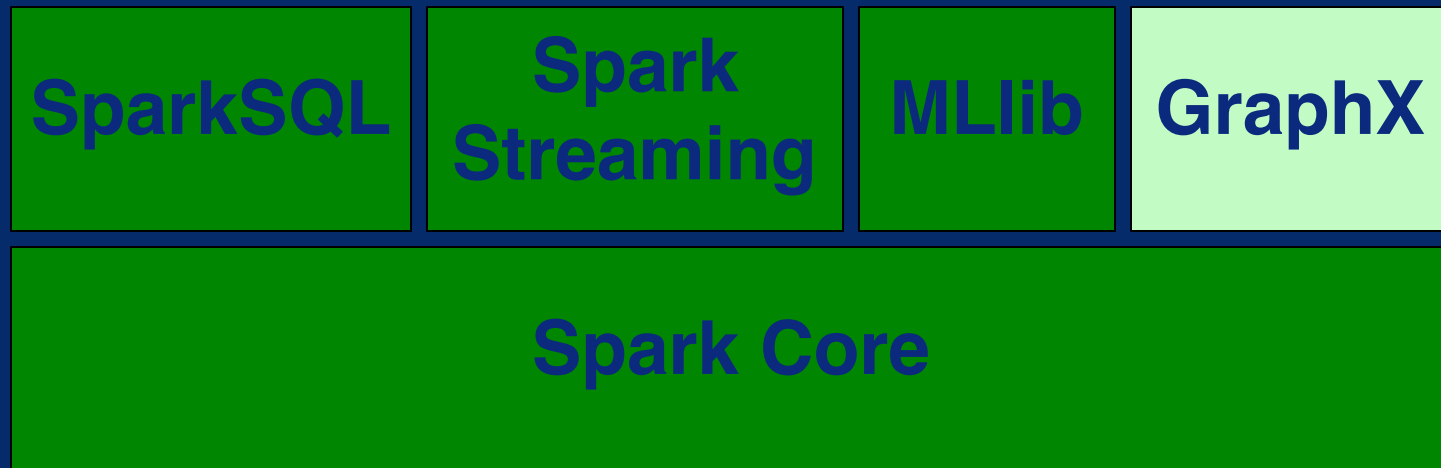
**Streaming processing**

# The Spark Stack



**Machine learning**

# The Spark Stack



**Graph analytics**

# The Spark Stack



**SparkSQL**

**Spark  
Streaming**

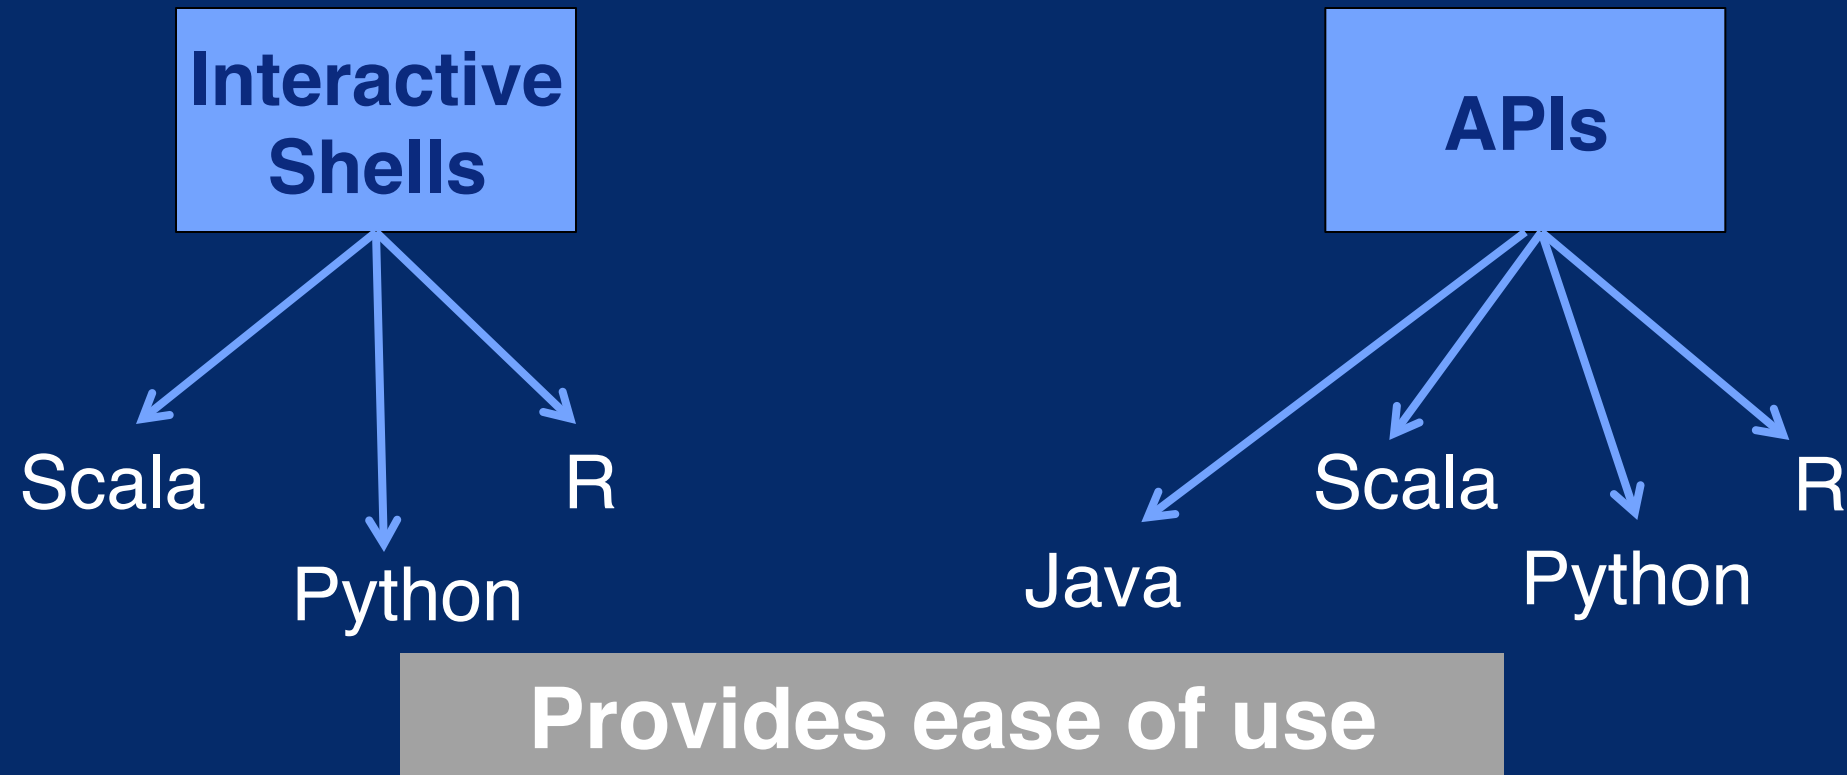
**MLlib**

**GraphX**

**Spark Core**

**Supports diverse analytics applications**

# Spark Interface




# In Memory Processing


Provides speed

**What does in memory processing mean?**






# MapReduce



# Resilient Distributed Datasets



# Resilient Distributed **Datasets**

*Dataset*

*Data storage created from:  
HDFS, S3, HBase, JSON, text,  
Local hierarchy of folders*

*Or created transforming  
another RDD*

# Resilient **Distributed** Datasets

*Distributed*

*Distributed across the cluster  
of machines*

*Divided in partitions, atomic  
chunks of data*

# Resilient Distributed Datasets

*Resilient*

*Recover from errors, e.g.  
node failure, slow processes*

*Track history of each  
partition, re-run*

# DataFrames & DataSets

DataFrame

DataSet

- **Extensions to RDDs**
- **Provide higher-level abstractions, improved performance, better scalability**

# Programming in Spark



# Creating RDDs

*Driver  
Program*

```
In [1]: lines = sc.textFile("hdfs://user/cloudera/words.txt")
```

# Creating RDDs

## *Driver Program*

```
In [1]: lines = sc.textFile("hdfs://user/cloudera/words.txt")
```

```
lines = sc.parallelize(["big", "data"])
```

# Creating RDDs

*Driver  
Program*

```
In [1]: lines = sc.textFile("hdfs://user/cloudera/words.txt")
```

```
lines = sc.parallelize(["big", "data"])
```

```
numbers = sc.parallelize(range(10), 3)
```

# Creating RDDs

*Driver  
Program*

```
In [1]: lines = sc.textFile("hdfs://user/cloudera/words.txt")
```

```
lines = sc.parallelize(["big", "data"])
```

```
numbers = sc.parallelize(range(10), 3)
```



[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]

# Creating RDDs

## Driver Program

```
In [1]: lines = sc.textFile("hdfs://user/cloudera/words.txt")
```

```
lines = sc.parallelize(["big", "data"])
```

```
numbers = sc.parallelize(range(10), 3)
```

[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]

Parallelize  
range output  
into 3 partitions

[0, 1, 2], [3, 4, 5], [6, 7, 8, 9]

# Creating RDDs

*Driver  
Program*

```
In [1]: lines = sc.textFile("hdfs://user/cloudera/words.txt")
```

```
lines = sc.parallelize(["big", "data"])
```

```
numbers = sc.parallelize(range(10), 3)
```

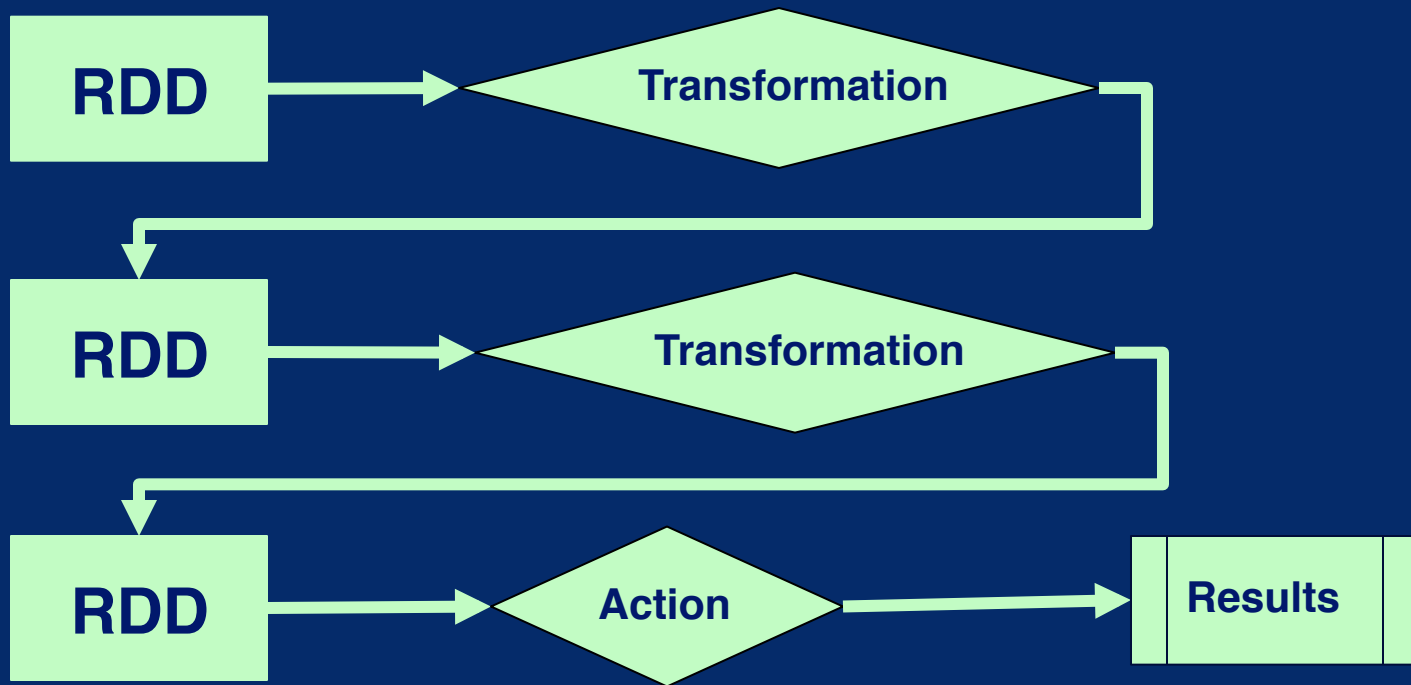
Parallelize  
range output  
into 3 partitions

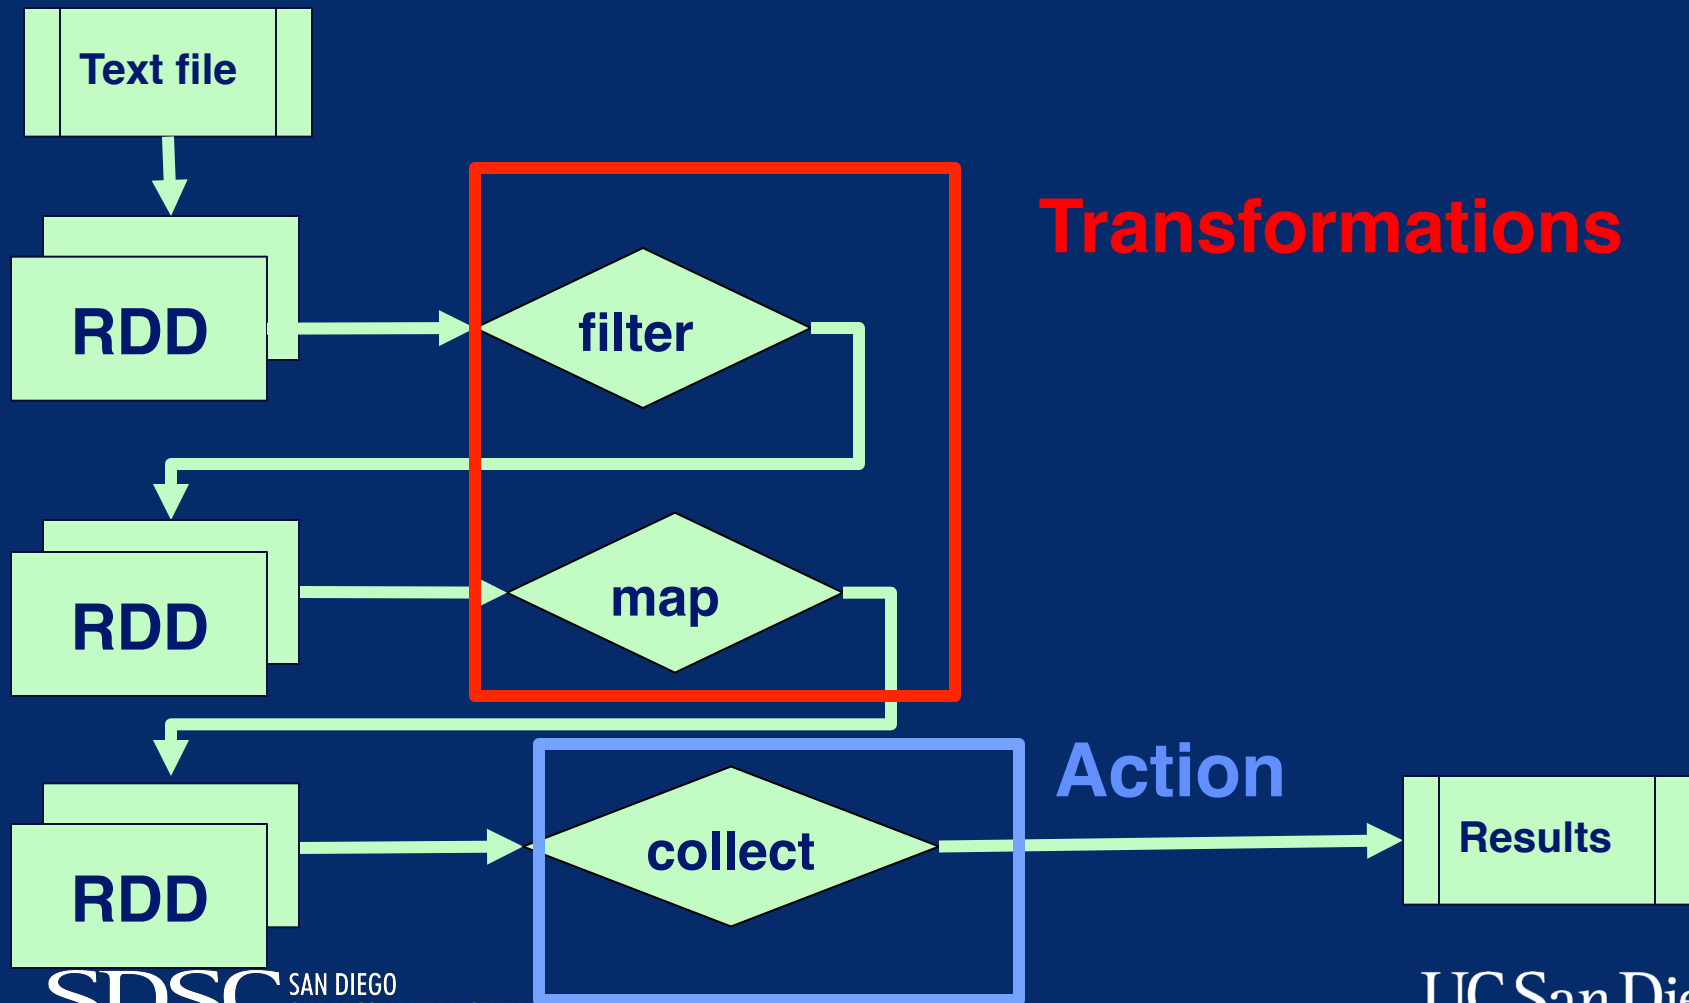
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]

```
numbers.collect()
```

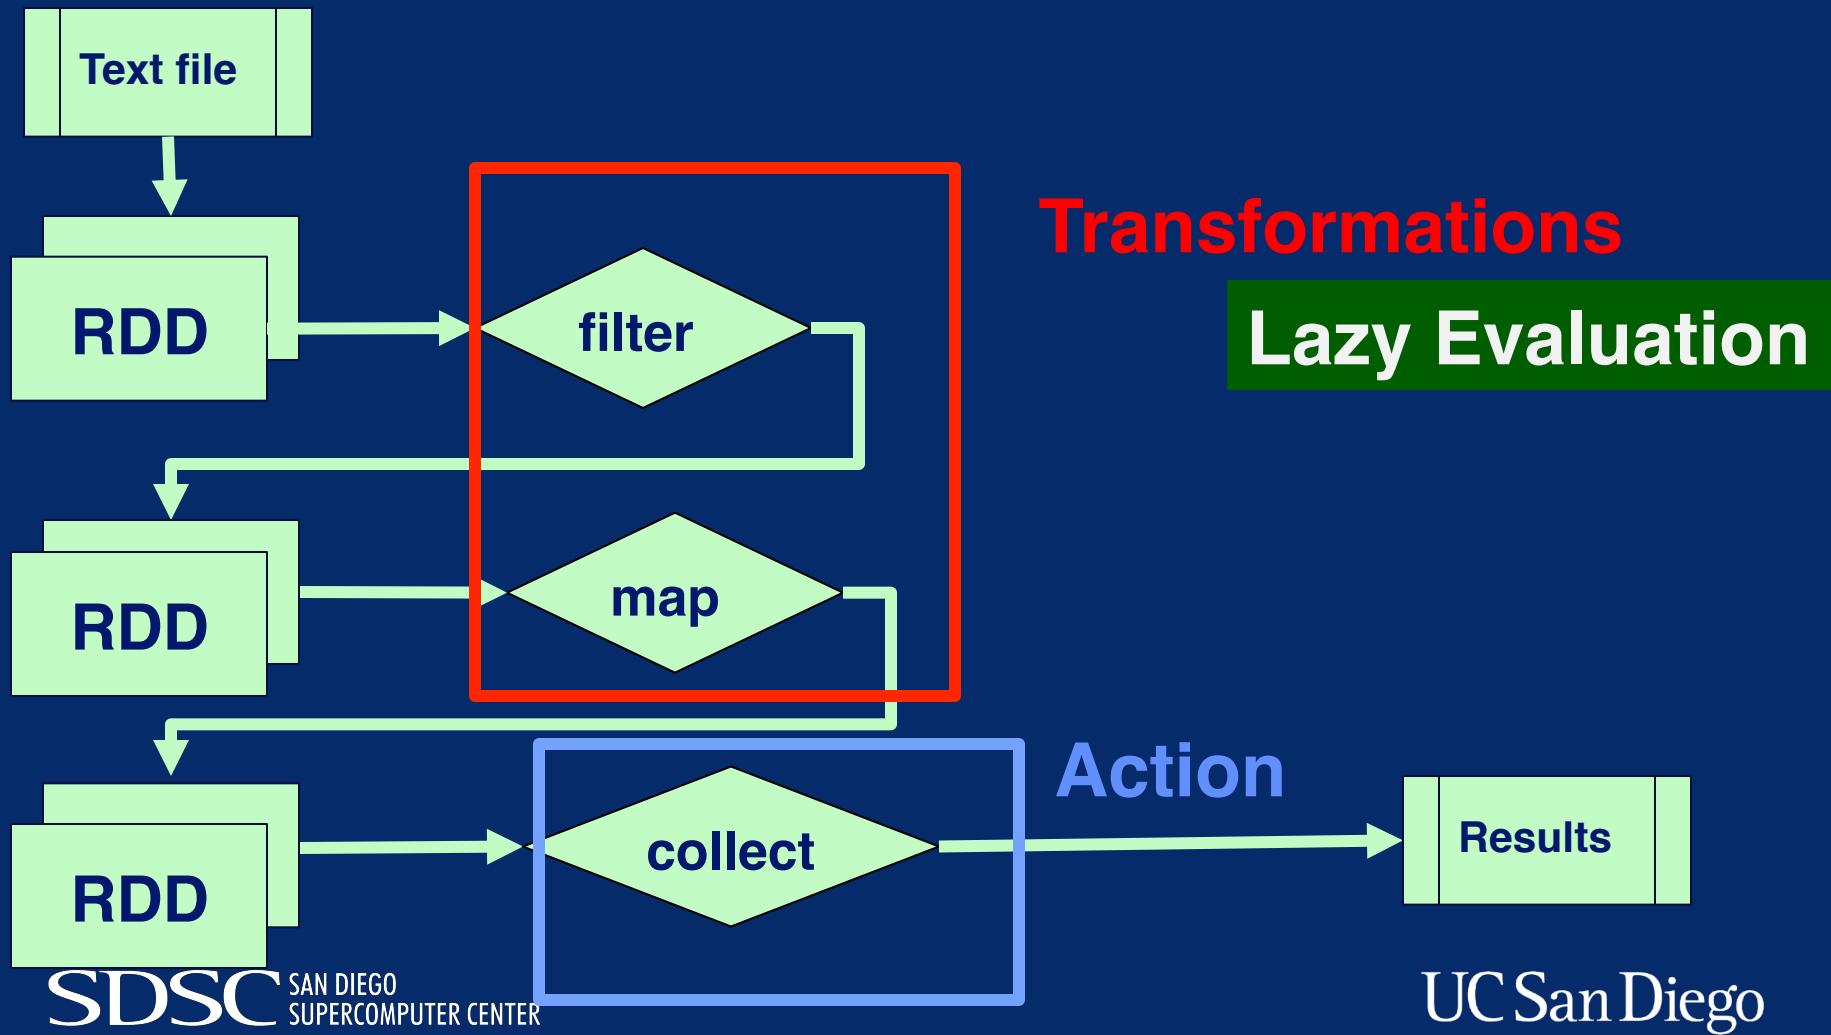
[0, 1, 2], [3, 4, 5], [6, 7, 8, 9]

# Processing RDDs









# Transformations & Actions

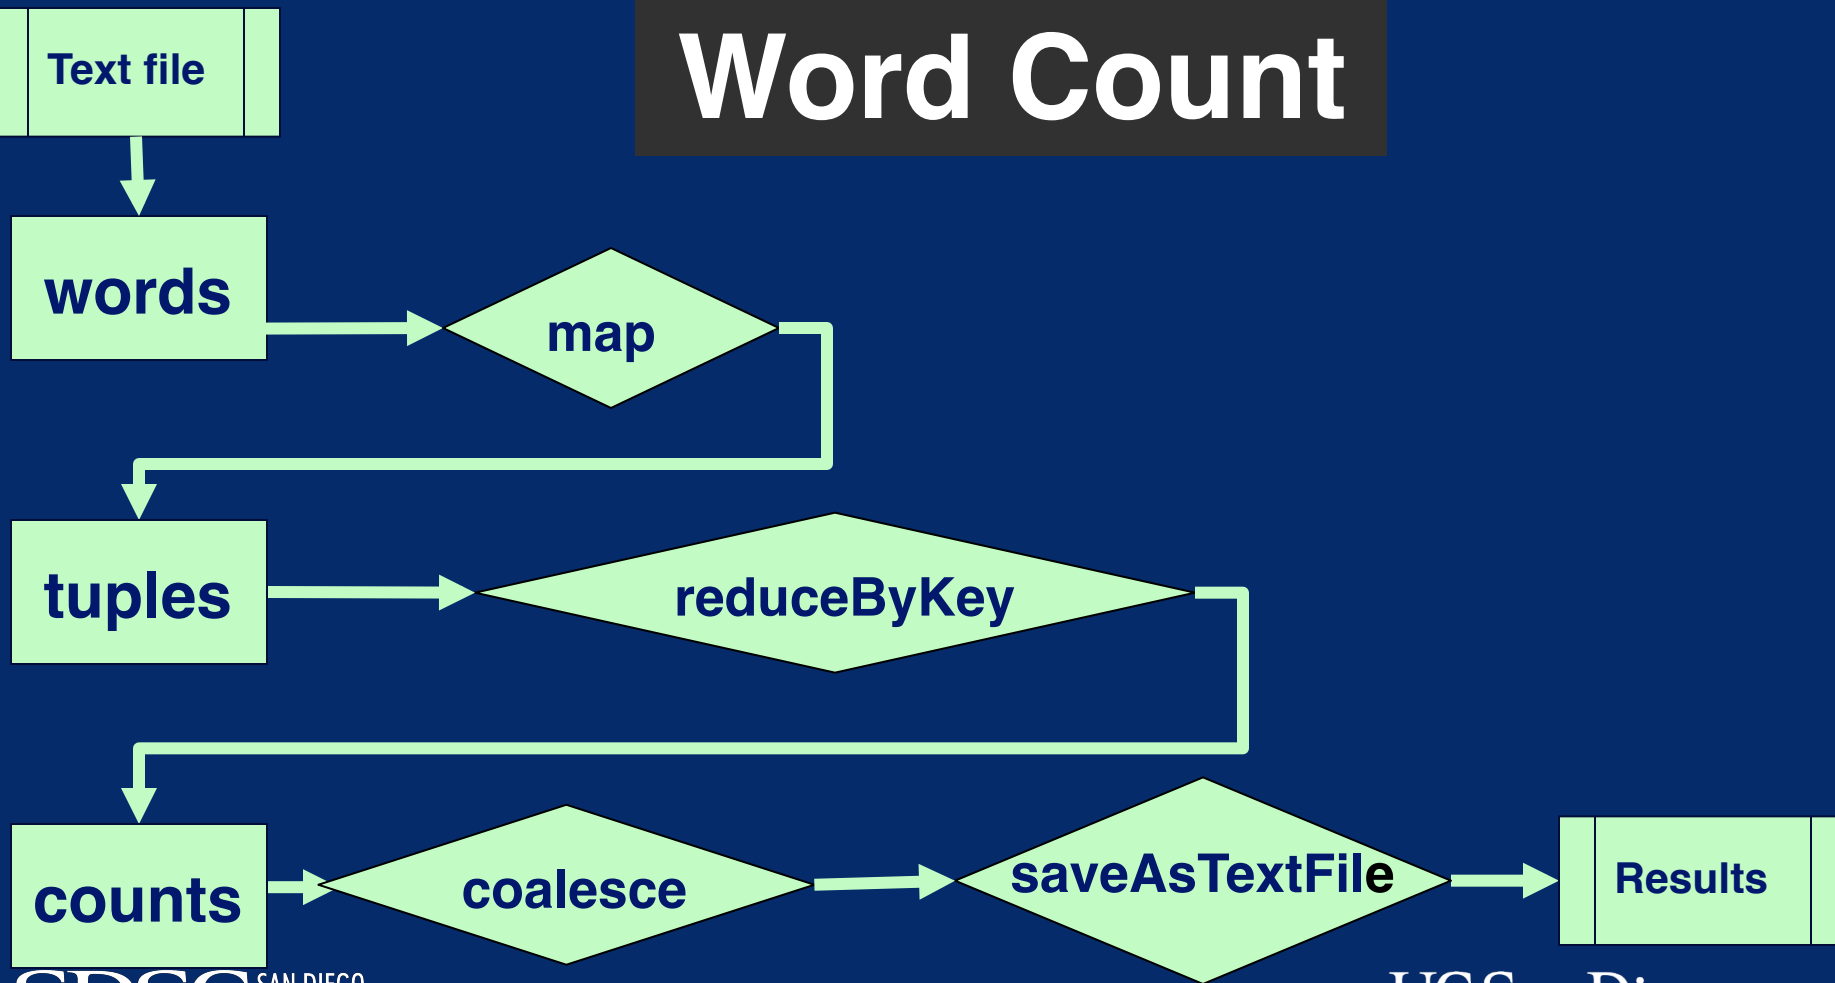
## Transformations

- **map**
- **filter**
- **coalesce**
- **reduceByKey**

## Actions

- **collect**
- **take**
- **reduce**
- **saveAsText**

# Word Count



# Programming in Spark

**Create RDDs**

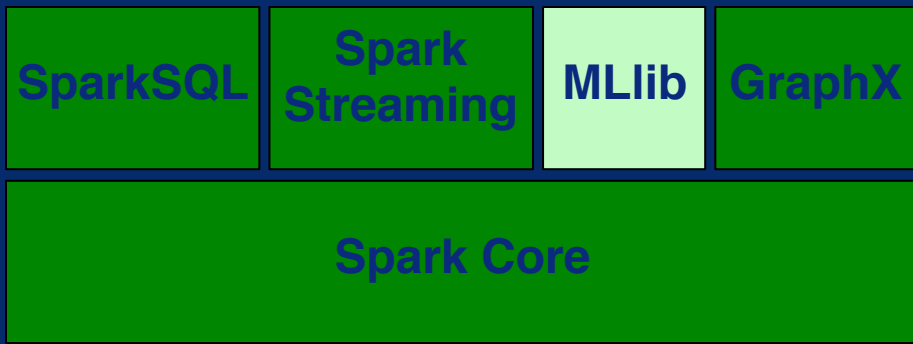


**Apply transformations**



**Perform actions**

# Spark MLlib: Machine Learning



# Spark MLlib

- Scalable machine learning library
- Provides distributed implementations of common machine learning algorithms and utilities
- Has APIs for Scala, Java, Python, and R

# MLlib Algorithms & Techniques

- Machine Learning
  - Classification, regression, clustering, etc.
  - Evaluation metrics
- Statistics
  - Summary statistics, sampling, etc.
- Utilities
  - Dimensionality reduction, transformation, etc.

# MLlib Example – Summary Statistics

- Compute column summary statistics

```
from pyspark.mllib.stat import Statistics
```

1

*# Data as RDD of Vectors*

```
dataMatrix = sc.parallelize([[1, 2, 3], [4, 5, 6], [7, 8, 9], [10, 11, 12]])
```

2

*# Compute column summary statistics.*

```
summary = Statistics.colStats(dataMatrix)
```

3

```
print(summary.mean())
```

4

```
print(summary.variance())
```

```
print(summary.numNonzeros())
```



# MLlib Example – Clustering

- Build k-means model for clustering

```
from pyspark.mllib.clustering import KMeans, KMeansModel  
from numpy import array
```

*# Read and parse data*

```
data = sc.textFile("data.txt")  
parsedData = data.map(lambda line:  
    array([float(x) for x in line.split(' ')]))
```

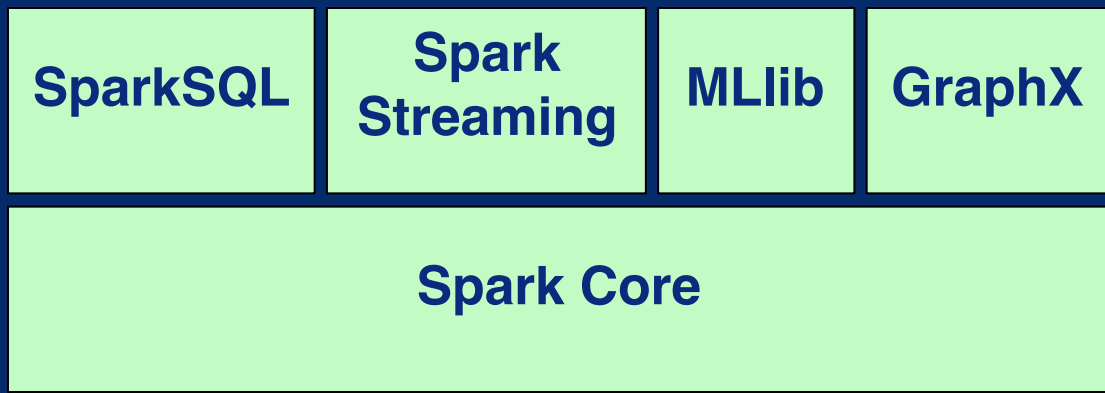
*# k-means model for clustering*

```
clusters = Kmeans.train (parsedData, k=3)
```

```
print(clusters.centers)
```

# Spark MLlib

- MLlib is Spark's machine learning library.
  - Distributed implementations
- Main categories of algorithms and techniques:
  - Machine learning
  - Statistics
  - Utilities for data preparation



- Spark core provides distributed computing
- Libraries support multiple analytics applications and workloads
- RDDs provide data parallelism & fault-tolerance
- MLlib provides scalable machine learning

# Q&A

